

Product Context — Wealth Management Platform

Living reference doc. Co-maintained by Ayush (PM) and Claude (PM partner). Update this file rather than re-explaining context in chat.

1. What We're Building

A multi-asset wealth management / portfolio tracking platform for the Indian market — positioned as a **deeper, better version of Mprofit**, not a clone. Phase 1 is **web-only** (no desktop/mobile/iOS yet).

Named competitors to date: Mprofit, Kuvera, PowerUp (Money), "ngen" (likely INDmoney or Fisdom — needs confirmation), passbook.ai, "Famili" (likely Fyers/a family-office tool — needs confirmation). More to come from Ayush.

2. Status — What's Actually Built

- **MF import via CAS parsing:**  Implemented and working, built from the Claude Code prompt (casparser-based PDF parsing → normalization → NAV enrichment via mfapi.in → XIRR/FIFO gains). This is live and in use now.
- **MF Central direct API/QR flow:** Not yet — gated behind AMFI ARN / SEBI RIA registration. **A registered partner is confirmed and in progress** for this. Until that's live (~5–10 months), CAS parsing remains the primary MF import path.
- **Equity/broker data:** Not started. Planned APIs: Zerodha, Groww, HDFC Securities, ICICI Direct, Angel One, Kotak Securities, Upstox, Axis Direct. Also considering NSDL/CDSL statements for cost-basis/market-value reconciliation, and free NSE/BSE data as a fallback for pricing.
- **Dashboard data feeds:** Google Finance API and Yahoo Finance API being evaluated for real-time dashboard pricing (scope TBD — which widgets need this vs. broker/NAV data).

3. Product Structure & Differentiators (vs. Mprofit)

Area	Mprofit's approach	Our planned approach
Signup/onboarding	Standard form-based signup	Questionnaire-style onboarding — more

Area	Mprofit's approach	Our planned approach
		user-centric, personalizes the experience from step 1
Account structure	Per-portfolio	One login → one family → up to 5 members (future phase)
MF import	MFCentral OTP/QR sync	CAS-PDF parsing now → MFCentral API via ARN partner later
Contract notes	Manual download + upload	Email automation: auto-fetch contract note emails, extract PDF password, feed to LLM+OCR pipeline — no manual download/upload step
Corporate actions	Automated application	Same, plus news-driven auto-update: monitor top-5 holdings per user for corporate action news (primarily via Moneycontrol), push via WhatsApp + email alerts
Analytics dashboard	Standard reporting	Reference: NGIN research website (Ayush to share specifics) + additional modules: brokerage analysis, STT analysis, GST analysis — advanced but customer-facing, not just raw numbers
Financial calculators	None visible	Calculators referencing Finology's toolset, rebuilt in a gamified dashboard format
FD / emergency fund tracking	Present	To be designed — approach TBD

4. Known Hard Problems (flagged, not yet solved)

1. **Data duplication:** if a mutual fund appears in both the MFCentral feed and a broker API feed (e.g., MF held via Zerodha Coin + also visible via MFCentral), we need dedup logic before this becomes a real integration — not an afterthought.
2. **Equity data sourcing:** no single free/reliable NSE/BSE API confirmed yet. Needs research: official NSE/BSE data licensing costs vs. broker-API-only coverage vs. scraping-adjacent free tiers (approach must stay within the same legitimacy bar we've held for MF data — no scraping/reverse-engineering of non-public endpoints).
3. **FD / emergency fund / cash-equivalent tracking:** no data source identified yet — likely manual entry only, at least in phase 1.
4. **Contract note email automation:** requires secure email access (OAuth to user's inbox), PDF password extraction/handling, and an LLM+OCR pipeline — meaningfully more complex and higher-trust than CAS parsing since it touches a live inbox connection.

5. Technical Direction (confirmed so far)

- **Phase 1 scope:** Web app only.
- **Storage:** AWS S3 (object storage) + Athena (query layer over S3) — this is a data-lake-style pattern, not a traditional RDBMS-first approach. Needs research: is Athena the right query layer for transactional portfolio data (which needs frequent point reads/writes), or is it better suited only for analytics/reporting workloads sitting alongside a proper OLTP database (Postgres) for the transactional core? This needs an explicit ADR before more is built on it.
- **MF import:** casparser (Python) pipeline — already built and working.
- **Money math:** Decimal-safe throughout (already a stated requirement from the earlier prompt).
- Broader tech stack (backend language/framework, frontend framework, hosting beyond S3/Athena) — **not yet decided**, flagged as needing dedicated research.

6. Team & Working Model

- Ayush is PM. Claude is being brought on as a **co-PM partner** for this project — meaning: helping structure the product (PRDs, prioritization, competitive analysis, opportunity mapping), not just executing tickets.
- Ayush has ~68 PM skills available (pm-skills plugin + Anthropic's product-management plugin). Skill usage plan is in Section 8 below.
- A separate engineering hire (first Software Engineer, genericized JD) is in progress in parallel — from the memory record, this thread is tracked separately.
- Target: **August build-out** of the core product (MF module refinement + equity module kickoff), per Ayush's stated timeline.

7. Open Questions — Need From Ayush Before Deeper Work Starts

Strategic

1. Full competitor list (only 5 named so far, more promised) — and are "ngen" and "Famili" the correct product names? (Couldn't confirm via search — possible transcription artifacts, same pattern as "Wealth Engine/PowerUp/Safebox" earlier.)
2. Target customer for phase 1 launch: retail DIY investor, HNI, or advisor/CA managing client portfolios? (Affects the questionnaire onboarding design and which broker APIs to prioritize.)
3. NGIN research website reference for dashboards — please share the actual link/screenshots; couldn't independently identify this source.
4. Finology calculators — which specific calculators (SIP, XIRR, retirement, tax) should be prioritized for the gamified rebuild?

Data/Integration 5. Which broker API access do we actually have or can realistically get for August (Kite Connect requires a developer account + fees; most others need partnership agreements)? This determines real scope vs. wishlist. 6. Status/timeline of the AMFI ARN partner relationship — even a rough date range for MFCentral API readiness helps sequence the roadmap. 7. Any existing NSE/BSE data vendor relationships, or starting from zero?

Process 8. Should I treat the "hiring a developer" thread and this thread as one merged context going forward, or keep them separate? 9. Confirm: do you want formal PRDs per module (MF, equity, dashboard, calculators) before engineering starts, or lighter specs given the Aug timeline pressure?

8. PM-Skill Usage Plan (proposed, sequenced)

I'm proposing a sequence — not using all 68 at once. Each stage below produces an artifact that feeds the next.

Stage A — Lock the strategic frame (do this before any more building)

- **discover-competitive-analysis** — proper structured teardown of Mprofit + Kuvera + PowerUp + others, once full competitor list is confirmed (upgrades our informal research into a real 2x2 + recommendations doc)
- **foundation-persona** — build 1–2 evidence-based personas (retail SIP investor, advisor-managing-clients) to stop us from designing in the abstract

- `define-jtbd-canvas` — the "job" investors are hiring us for (this shapes the questionnaire-onboarding design directly)
- `foundation-lean-canvas` or `tool-foundation-sprint-basics` — pick one: lean canvas is faster (single doc), the Foundation Sprint tool is heavier but forces explicit team alignment on target customer/differentiators/competitors. Given August pressure, I'd lean lean-canvas now, sprint later if misalignment shows up.

Stage B — Turn strategy into scoped work

- `define-opportunity-tree` — connects the outcome ("investor sees true consolidated wealth") to real customer opportunities (data duplication pain, manual contract notes, etc.) before jumping to solutions — this directly prevents the solution-first jumps this conversation has already been prone to (e.g., jumping to "email automation" before confirming it's the highest-value opportunity)
- `define-prioritization-framework` (RICE or ICE) — score MF-polish vs. equity-integration vs. dashboard vs. calculators for the August window
- `develop-adr` — **needed now** for the S3+Athena storage question in Section 5; this is a real architectural fork, not a detail
- `deliver-prd` — one per module once prioritized (MF v2, Equity v1, Analytics Dashboard v1)
- `deliver-user-stories` + `deliver-acceptance-criteria` — engineering handoff once PRDs exist
- `deliver-edge-cases` — specifically for the data-duplication problem and CAS-parsing failure modes

Stage C — Execution support (ongoing, as the eng hire ramps)

- `foundation-meeting-agenda` / `foundation-meeting-recap` — standups/planning once the engineer is onboarded
- `product-management:sprint-planning` — August sprint scoping against real capacity
- `measure-instrumentation-spec` — before shipping the dashboard, define what we're tracking
- `utility-pm-critic` — adversarial review pass on PRDs before they go to engineering (cheap insurance)

Stage D — Post-launch

- `iterate-retrospective`, `iterate-lessons-log`, `measure-experiment-results` (e.g., A/B testing the questionnaire onboarding vs. a control), `foundation-okr-writer` for the next quarter

My recommendation for right now: Stage A first — specifically `discover-competitive-analysis` (once you send the fuller competitor list) and

`define-jtbd-canvas`. Everything downstream gets sharper if those two are solid, and they're fast.

9. What I Still Need From You Overall

- Answers to Section 7's open questions (even partial/rough is fine — I'll flag remaining gaps as we go)
- The actual roadmap prompt you mentioned is coming next
- Full competitor list when ready
- Any existing PRD/spec docs, design references, or research files you want loaded into project context